

# CS 474/574 Machine Learning

## 1. Introduction

Prof. Dr. Forrest Sheng Bao  
Dept. of Computer Science  
Iowa State University  
Ames, IA, USA

January 27, 2021

# Why Machine Learning (ML)

- ▶ How computers know how to do things?

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
  2. learning: feeding the computer many bilingual documents



# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
  2. learning: feeding the computer many bilingual documents
- ▶ Example 2: sorting

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
  2. learning: feeding the computer many bilingual documents
- ▶ Example 2: sorting
  1. programming: Quicksort, etc.

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
  2. learning: feeding the computer many bilingual documents
- ▶ Example 2: sorting
  1. programming: Quicksort, etc.
  2. learning: feeding the computer many pairs of unsorted and sorted list of numbers.

# Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
  1. programming: steps detailed by human programmer
  2. learning: without being specifically told
- ▶ Example 1: machine translation
  1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
  2. learning: feeding the computer many bilingual documents
- ▶ Example 2: sorting
  1. programming: Quicksort, etc.
  2. learning: feeding the computer many pairs of unsorted and sorted list of numbers.
- ▶ The first approach in the context of AI is also called rule-based system or expert system, e.g. MyCin, Grammarly.

# Why ML is attractive

- ▶ We are lazy. We want to shift the heavy lifting to the computers.

# Why ML is attractive

- ▶ We are lazy. We want to shift the heavy lifting to the computers.
- ▶ We are incompetent. No kidding! Sometimes it is very difficult to come up with step-by-step instructions.

# Why ML is attractive

- ▶ We are lazy. We want to shift the heavy lifting to the computers.
- ▶ We are incompetent. No kidding! Sometimes it is very difficult to come up with step-by-step instructions.
- ▶ Examples: Self-driving, AlphaGo, Automated circuit routing, Machine translation, Commonsense reasoning, text entailment, Document generation, auto-reply of messages/emails, fly a helicopter inversely, van-Gogh-lize paints.

# Why ML is attractive

- ▶ We are lazy. We want to shift the heavy lifting to the computers.
- ▶ We are incompetent. No kidding! Sometimes it is very difficult to come up with step-by-step instructions.
- ▶ Examples: Self-driving, AlphaGo, Automated circuit routing, Machine translation, Commonsense reasoning, text entailment, Document generation, auto-reply of messages/emails, fly a helicopter inversely, van-Gogh-lize paints.
- ▶ It is a dream. “Creating an artificial being has been the dream since the beginning of science.” – Movie A.I., Spielberg et al., 2001



# Functions

- ▶ A function is a many(incl. one)-to-one mapping between data!

# Functions

- ▶ A function is a many(incl. one)-to-one mapping between data!
- ▶ It does not need to have an analytic form and the domain/range does not have to be of real numbers.

# Functions

- ▶ A function is a many(incl. one)-to-one mapping between data!
- ▶ It does not need to have an analytic form and the domain/range does not have to be of real numbers.
- ▶ An example of edge detection in Wolfram Mathematica:

In[1]:=

EdgeDetect[



]

Out[1]:=



# Functions

- ▶ A function is a many(incl. one)-to-one mapping between data!
- ▶ It does not need to have an analytic form and the domain/range does not have to be of real numbers.
- ▶ An example of edge detection in Wolfram Mathematica:

In[1]:=

EdgeDetect[



]

Out[1]:=



- ▶ Another example in machine translation

# Supervised ML as function approximation

## Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.

## Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .

# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .



# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .
  - ▶ **Supervised** because we provide many pairs of  $x$ 's and  $y$ 's for the computer to know the difference between  $\hat{y}$  and  $y$  on a large pool of samples.

# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .
  - ▶ **Supervised** because we provide many pairs of  $x$ 's and  $y$ 's for the computer to know the difference between  $\hat{y}$  and  $y$  on a large pool of samples.
  - ▶ Examples: object detection from images, Flavia, CPU branch prediction, COVID-19 diagnosis from blood profile, Epileptic EEG recognition, depression treatment from brain shapes.

# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .
  - ▶ **Supervised** because we provide many pairs of  $x$ 's and  $y$ 's for the computer to know the difference between  $\hat{y}$  and  $y$  on a large pool of samples.
  - ▶ Examples: object detection from images, Flavia, CPU branch prediction, COVID-19 diagnosis from blood profile, Epileptic EEG recognition, depression treatment from brain shapes.
  - ▶ Beyond categorization/classification: Mflux, Review helpfulness prediction, Document summarization, predict house prices

# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .
  - ▶ **Supervised** because we provide many pairs of  $x$ 's and  $y$ 's for the computer to know the difference between  $\hat{y}$  and  $y$  on a large pool of samples.
  - ▶ Examples: object detection from images, Flavia, CPU branch prediction, COVID-19 diagnosis from blood profile, Epileptic EEG recognition, depression treatment from brain shapes.
  - ▶ Beyond categorization/classification: Mflux, Review helpfulness prediction, Document summarization, predict house prices
- ▶ Unsupervised, finding  $\hat{f}(x)$  without ground truth

# Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding  $\hat{f}(x) \approx f(x)$  with ground truth provided by human.
  - ▶ Let  $x$  and  $y$  be two (vectors of) variables, and a function connecting them  $y = f(x)$  But only god knows  $f$ .
  - ▶ We construct another function  $\hat{f}$  to approximate  $f$  such that  $\hat{y} = \hat{f}(x) \approx y = f(x)$  for a(ny) given  $x$ .
  - ▶ **Supervised** because we provide many pairs of  $x$ 's and  $y$ 's for the computer to know the difference between  $\hat{y}$  and  $y$  on a large pool of samples.
  - ▶ Examples: object detection from images, Flavia, CPU branch prediction, COVID-19 diagnosis from blood profile, Epileptic EEG recognition, depression treatment from brain shapes.
  - ▶ Beyond categorization/classification: Mflux, Review helpfulness prediction, Document summarization, predict house prices
- ▶ Unsupervised, finding  $\hat{f}(x)$  without ground truth
- ▶ Reinforcement, let the machine find ground truth itself

# Representation of $x$

- ▶  $x$  is usually not a simple (vector of) number(s). How to tell it to a computer?

# Representation of $x$

- ▶  $x$  is usually not a simple (vector of) number(s). How to tell it to a computer?
- ▶ Example: bananas vs. apples

# Representation of $x$

- ▶  $x$  is usually not a simple (vector of) number(s). How to tell it to a computer?
- ▶ Example: bananas vs. apples
- ▶ **Feature engineering**: manually craft functions to **extract** features from raw data, e.g., SIFT, bag-of-words.



# Representation of $x$

- ▶  $x$  is usually not a simple (vector of) number(s). How to tell it to a computer?
- ▶ Example: bananas vs. apples
- ▶ **Feature engineering**: manually craft functions to **extract** features from raw data, e.g., SIFT, bag-of-words.
- ▶ Automated feature extraction in deep learning: E.g., filters in CNNs.

# Representation of $x$

- ▶  $x$  is usually not a simple (vector of) number(s). How to tell it to a computer?
- ▶ Example: bananas vs. apples
- ▶ **Feature engineering**: manually craft functions to **extract** features from raw data, e.g., SIFT, bag-of-words.
- ▶ Automated feature extraction in deep learning: E.g., filters in CNNs.
- ▶ If  $x$  involves categorical values (e.g., gender), there are usually two approaches: **One-hot encoding** and **embedding** (in DL context, to be discussed later).

## Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  
 $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  
 $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  
 $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .
- ▶ The process of finding the approximation function  $\hat{f}$  is called **training** or **learning**.

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .
- ▶ The process of finding the approximation function  $\hat{f}$  is called **training** or **learning**.
- ▶  $\hat{f}$  is called a **model** or an **estimator**.



# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .
- ▶ The process of finding the approximation function  $\hat{f}$  is called **training** or **learning**.
- ▶  $\hat{f}$  is called a **model** or an **estimator**.
- ▶  $\mathbf{X}_i$ : an **input** (especially when raw data is used as the input) or **feature vector** (if using feature engineering).

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .
- ▶ The process of finding the approximation function  $\hat{f}$  is called **training** or **learning**.
- ▶  $\hat{f}$  is called a **model** or an **estimator**.
- ▶  $\mathbf{X}_i$ : an **input** (especially when raw data is used as the input) or **feature vector** (if using feature engineering).
- ▶  $\mathbf{y}_i$ , often  $\in \mathbb{R}^1$  a **label** (in classification) or **target** (used more generally and lately).

# Supervised ML

- ▶ Given many pairs of inputs and outputs:  
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  such that  $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function  $\hat{f}$  that approximates the function  $f$ .
- ▶ “approximate”: usually  $\min ||\hat{f}(x) - f(x)||^p$  where  $p$  is usually 1 or 2. See  $\ell_p$ -norm .
- ▶ The process of finding the approximation function  $\hat{f}$  is called **training** or **learning**.
- ▶  $\hat{f}$  is called a **model** or an **estimator**.
- ▶  $\mathbf{X}_i$ : an **input** (especially when raw data is used as the input) or **feature vector** (if using feature engineering).
- ▶  $\mathbf{y}_i$ , often  $\in \mathbb{R}^1$  a **label** (in classification) or **target** (used more generally and lately).
- ▶ Classification vs. Regression: When  $y$  is continuous or discrete. In modern DL context, such division is usually no mentioned, especially in generative tasks.